

Introduction to Artificial Intelligence

LIAW, RUNG-TZUO

Department of Computer Science and Information Engineering

Fu Jen Catholic University, Taiwan

Chapter 3

Solving Problems by Searching

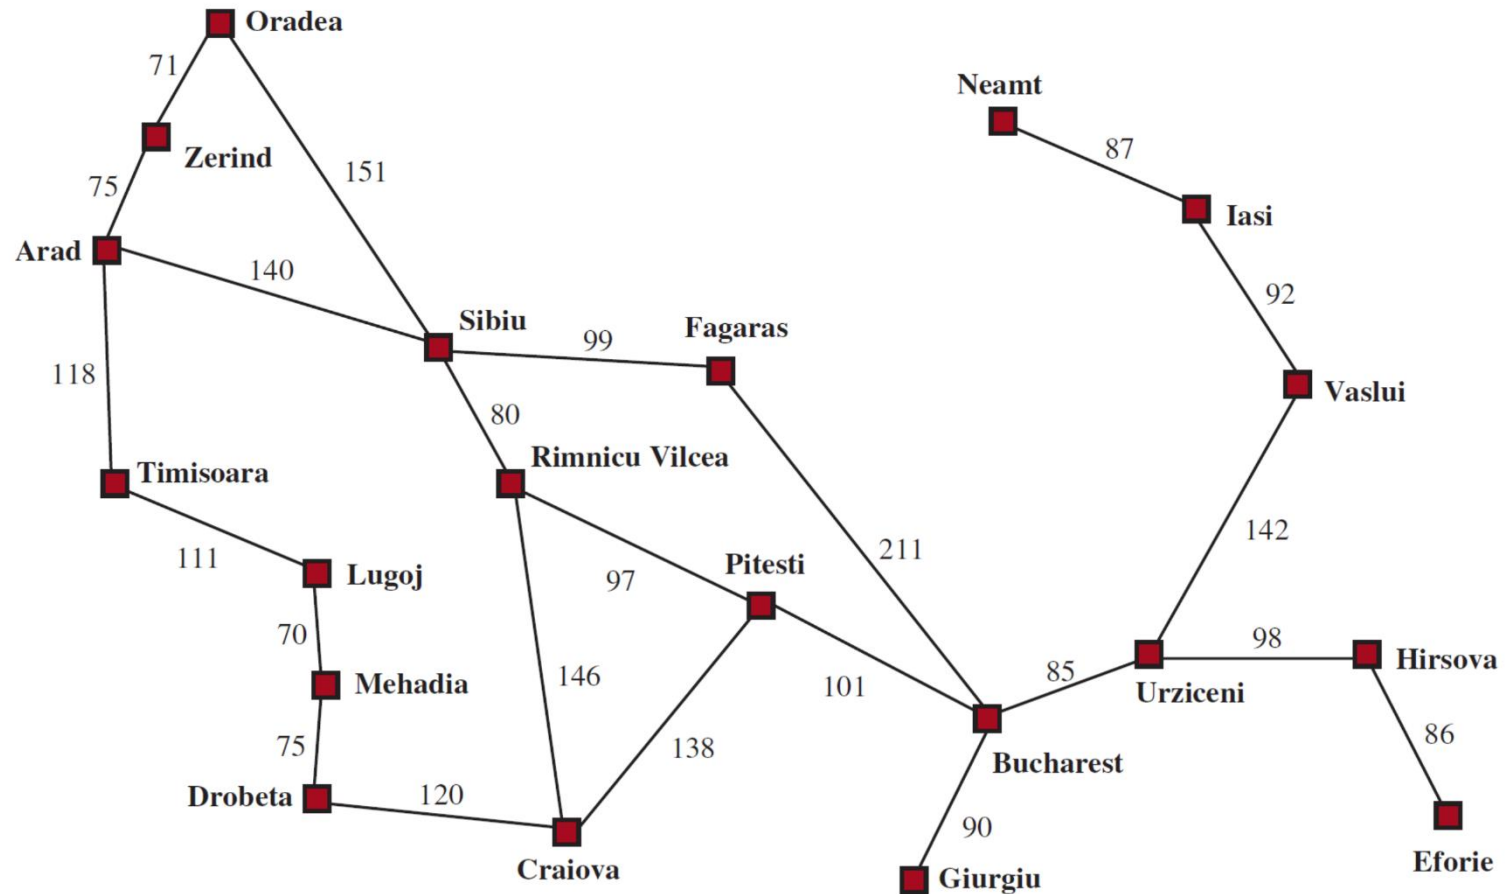
Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Example: Romania



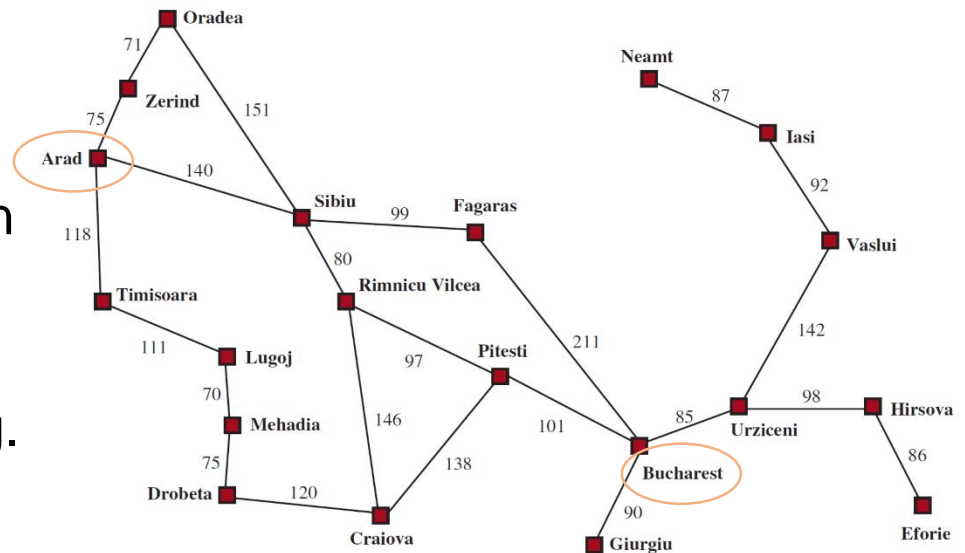
A simplified road map of part of Romania, with road distances in miles.

Example: Romania

- **On holiday in Romania**

Currently in Arad; flight leaves from Bucharest tomorrow

- Formulate goal
 - **Be in Bucharest**
- Formulate problem
 - States: various cities
 - Actions: drive between cities
- Find solution
 - Sequence of cities e.g. Arad, Sibiu, Fagaras, Bucharest, ...



Problem-Solving Agent

- A kind of **goal-based agent**
 - Goal: solve the problem
- **Four general steps in problem solving:**
 - **Goal** formulation
 - What are the successful world states
 - **Problem** formulation
 - What actions and states to consider given the goal
 - **Search**
 - Determine the possible *sequence of actions* that lead to the states of known values and then choosing the best sequence
 - **Execute**
 - Give the solution perform the actions

Problem-Solving Agent

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*)
returns an action

static: *seq*, an action sequence

state, a description of the current world state

goal, a goal

problem, a problem formulation

state \leftarrow UPDATE-STATE (*state*, *percept*)

if *seq* is empty **then do**

goal \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state*, *goal*)

seq \leftarrow SEARCH(*problem*)

action \leftarrow FIRST(*seq*)

seq \leftarrow REST(*seq*)

return *action*

Problem Types

- **Deterministic, fully observable → single state problem**
 - Agent knows exactly which state it will be in
 - Solution is a sequence

Goal Formulation

- **Determine what success means**
- **Perhaps the most important fundamental question**

Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Problem Formulation

- **A problem is defined by:**

- Initial state, e.g. Arad
- Successor function $S(X)$ = set of <action, state> pairs
 - e.g. $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \langle \text{ , } \rangle, \dots \}$
 - Initial state + successor function = state space
- Goal test, can be
 - Explicit, e.g. $x = \text{'at Bucharest'}$
 - Implicit, e.g. $\text{checkmate}(x)$
- Path cost (additive)
 - e.g. sum of distances, number of actions executed, ...
 - $c(x, a, y)$ is the step cost, assumed to be ≥ 0


- A **solution** is a sequence of actions from initial to goal state.
- **Optimal solution** has the lowest path cost.

State Space


- **Real world is absurdly complex**
 - State space must be **abstracted** for problem solving
 - (Abstract) state = set of real states
 - (Abstract) action = complex combination of real actions
 - e.g. Arad → Zerind represents a complex set of possible routes, detours, rest stops, etc.
 - The abstraction is valid if the path between two states is reflected in the real world
 - Each abstract action should be “easier” than the real problem
 - (Abstract) solution = set of real paths that are solutions in the real world

Example: Vacuum World

- Consider an environment with two cells and one vacuum




Example: Vacuum World



- **Problem Formulation**

- States?
- Initial state?
- Actions?
- Goal test?
- Path cost?

Example: Vacuum World



- **Problem Formulation**

- States? (vacuum, dust)
- Initial state? Any
- Actions? $L = \text{Left}$, $R = \text{Right}$, $S = \text{Suck}$.
- Goal test? Check if no dust exists.
- Path cost? Number of actions to reach goal.

Example: 8-Puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- **Problem Formulation**

- States?
- Initial state?
- Actions?
- Goal test?
- Path cost?

Example: 8-Puzzle

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

- **Problem Formulation**


- States? Integer location of each tile. How many of them?
- Initial state? Any state
- Actions? (tile, direction)
where direction is one of *{Left, Right, Up, Down}*
- Goal test? Check whether goal configuration is reached
- Path cost? Number of actions to reach goal

Example: 8 queens problem



- Incremental vs. complete state formulation
 - Incremental formulation starts with an empty state and involves operators that augment the state description
 - A complete state formulation starts with all 8 queens on the board and moves them around


8 queens problem: representation is key



- **Incremental Formulation**

- States? Any arrangement of 0 to 8 queens
- Initial state? No queens
- Actions? Add queen in empty square
- Goal test? 8 queens on board and none attacked
- Path cost? None but $64 \cdot 63 \cdot \dots \cdot 57 \sim 3E14$ states


8 queens problem: a better representation is key



- **Another Incremental Formulation**

- States? n queens on the board n left-most columns
- Initial state? No queens
- Actions? Add queen in leftmost empty column
- Goal test? 8 queens on board and none attacked
- Path cost? None but 2057 states

8 queens problem: a better representation is key



- **Complete-state Formulation**

- States? Position of 8 queens
- Initial state? Any possible positions
- Actions? Move a queen to another place
- Goal test? 8 queens on board and none attacked
- Path cost? #movements

n queens problem

- A solution is a goal node, not a path to this node (typical of design problem)
- Number of states in state space:
 - 8-queens $\rightarrow 2,057$
 - 100-queens $\rightarrow 10^{52}$
- But techniques exist to solve n-queens problems efficiently for large values of n



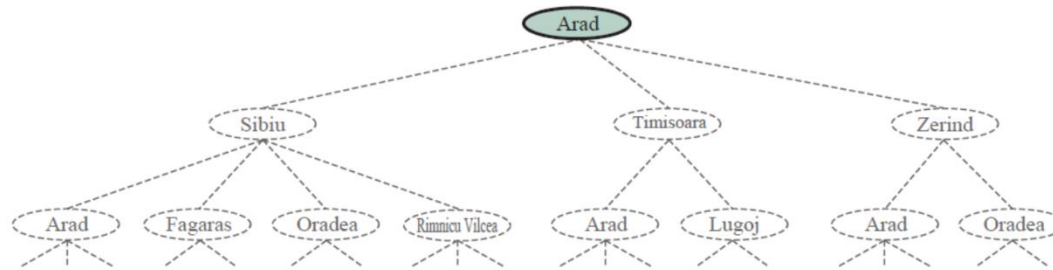
Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Basic Search Algorithms

- How do we **find the solutions** of previous problems?
 - **Search** the state space
 - Remember complexity of space depends on state representation
 - Here: search through **explicit tree** generation
 - ROOT= initial state
 - **Expand** current state
 - Generate new set of states (nodes and leafs) by applying *successor function* to current state
 - **Search strategy** chooses which state to expand
 - More generally, search generates a *graph*
 - Same state through multiple paths

Simple Tree Search Example



Simple Tree Search Example

function TREE-SEARCH(*problem*, *strategy*)

returns a solution or failure

initialize search tree using the initial state

loop do


if no candidates for expansion **return** failure

choose leaf node for expansion as per *strategy*

if node contains goal state **then return** solution

else expand the node and add resulting nodes
to the search tree


Simple Tree Search Example



A sequence of search trees generated by a graph search on the Romania problem

State Space vs. Search Tree

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure belong to a search tree



- A node has a parent, children, ... and includes path cost, depth, ...
- Here node = $\langle \text{state, parent-node, action, path-cost, depth} \rangle$
- FRINGE = contains generated nodes which are not yet expanded

Tree Search Algorithm

function TREE-SEARCH(problem, fringe) **returns** a solution or failure

fringe \leftarrow INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)

loop do

if EMPTY(fringe) **then return** failure

node \leftarrow REMOVE-FIRST(fringe)

if GOAL-TEST[problem] applied to STATE[node] succeeds

then return SOLUTION(node)

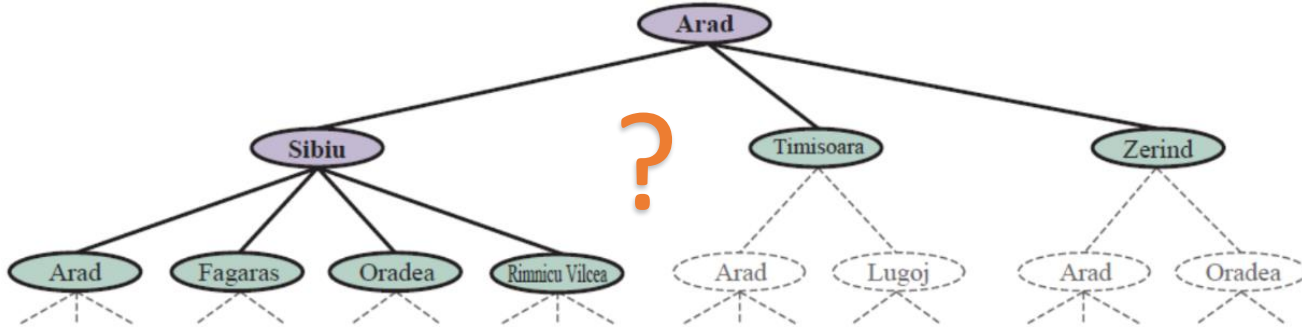
fringe \leftarrow INSERT-ALL(EXPAND(node, problem), fringe)

Tree Search Algorithm.

```
function EXPAND(node, problem) returns a set of nodes
    successors  $\leftarrow$  the empty set
    for each  $\langle \textit{action}, \textit{result} \rangle$  in SUCCESSORFN[problem](STATE[node]) do
        s  $\leftarrow$  a new NODE
        STATE[s]  $\leftarrow$  result
        PARENT-NODE[s]  $\leftarrow$  node
        ACTION[s]  $\leftarrow$  action
        PATH-COST[s]  $\leftarrow$  PATH-COST[node] +
                               STEP-COST(STATE[node], action, result)
        DEPTH[s]  $\leftarrow$  DEPTH[node] + 1
        add s to successors
    return successors
```

Search Strategies

- A **strategy** is defined by picking the order of **node expansion**
 - Should consider performance



Search Strategies (cont'd)

- **Problem-solving performance is measured in four ways:**
 - **Completeness**
 - Does it always find a solution if one exists?
 - **Optimality**
 - Does it always find the least-cost (optimal) solution?
 - (Is what it found always the least-cost (optimal) solution?)
 - **Time Complexity**
 - Number of nodes generated/expanded?
 - **Space Complexity**
 - Number of nodes stored in memory during search?

Search Strategies.

- **Time and space complexity are measured in terms of problem difficulty defined by:**
 - b : branching factor or maximum #successors of any node
 - d : depth of the shallowest goal node
 - m : maximum depth of the state space (may be ∞)

Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**


Uninformed Search Strategies

- **Uninformed search**

- a.k.a. blind search
- Uses only information available in problem definition
 - When strategies can determine whether one non-goal state is *better than* another → **informed** or **heuristic** search
- Categories defined by expansion algorithm:
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search
 - Bidirectional search

Breadth-First Search (BFS)

- **Expand shallowest unexpanded node**
 - Implementation: fringe is a FIFO queue
 - E.g. breadth-first search on a simple binary tree
 - At each stage, the node to be expanded next is indicated by the triangular marker



Breadth-First Search (BFS)

- **Completeness:**

- Does it always find a solution if one exists?
- YES
 - If shallowest goal node is at some finite depth d
 - Condition: If b is **finite** (maximum number of successive nodes is finite)

- **Optimality:**

- Does it always find the least-cost (optimal) solution?
- In general, YES
 - unless actions have different cost (i.e. path cost is a decreasing function)

Breadth-First Search (BFS)

- **Time complexity:**

- Assume a state space where every state has b successors
 - Root has b successors
 - each node at the next level has again b successors (total b^2)
 - \dots
 - Assume solution is at depth d
 - **Worst case:** expand all but the last node (goal) at depth d
 - Total #nodes generated:
$$b + b^2 + b^3 + \dots + b^d + (b^{d+1} - b) = O(b^{d+1})$$

- **Space complexity:**

- Idem $O(b^{d+1})$ if each node is retained in memory

Breadth-First Search (BFS)

- **Two lessons**

- **Memory** requirements are a bigger problem than its execution time
- **Exponential** complexity search problems cannot be solved by uninformed search for any but the smallest instances
 - e.g. $b=10$; 10,000 nodes/s; 1000 bytes/node

DEPTH	NODES	TIME	MEMORY
2	1100	0.11 seconds	1 megabyte
4	111100	11 seconds	106 megabytes
6	10^7	19 minutes	10 gigabytes
8	10^9	31 hours	1 terabyte
10	10^{11}	129 days	101 terabytes
12	10^{13}	35 years	10 petabytes
14	10^{15}	3523 years	1 exabyte

Uniform-Cost Search (UCS)

- **Extension of BFS**

- Expand node with lowest path cost
- UCS is the same as BFS when all step-costs are equal
- Implementation: Fringe = queue ordered by path cost

Uniform-Cost Search (UCS)

- **Completeness:**

- Does it always find a solution if one exists?
- YES, if step-cost $> \epsilon$ (small positive constant) (why?)
 - The search will get stuck in a infinite loop if a node has a zero-cost action leading back to the same state

- **Optimality:**

- Does it always find the least-cost (optimal) solution?
- YES, if complete
 - Nodes expanded in order of increasing path cost

Uniform-Cost Search (UCS)

- **Time complexity:**


- Assume C^* the cost of the optimal solution
- Assume that every action costs at least ε
- Worst-case: $O(b^{1+C^*/\varepsilon})$

- **Space complexity:**

- Idem to time complexity

Depth-First Search (DFS)

- **Expand deepest unexpanded node**
 - Implementation: fringe is a LIFO queue (=stack)
 - e.g.



A dozen steps in the progress of a DFS on a binary tree from start state A to goal M

Depth-First Search (DFS)

- **Completeness:**

- Does it always find a solution if one exists?
- **No**
 - e.g. the left subtree is unbounded depth but no solution
 - DFS is complete only if search space is finite and no loops

- **Optimality:**

- Does it always find the least-cost (optimal) solution?
- **No**
 - e.g. Assume nodes J and C contain goal states (But only C is optimal)

Depth-First Search (DFS)

- **Time complexity:** $O(b^m)$
 - Terrible if m (max. depth) is much larger than d (depth of optimal solution)
 - But if many solutions, then faster than BFS
- **Space complexity:** $O(bm)$
 - Backtracking search uses even **less memory**
 - One successor instead of all b
 - DFS' advantage over BFS: $O(b^{d+1})$

Depth-Limited Search

- **Limits the search depth of DFS to l**
 - Solves the infinite-path problem
 - i.e. nodes at depth l have no successors
 - If $l < d$ then incompleteness results
 - If $l > d$ still can be not optimal (why?)
 - Problem knowledge can be used
- **Time complexity:** $O(b^l)$
- **Space complexity:** $O(bl)$

Iterative Deepening Search (IDS)

- **A general strategy to find best depth limit /**
 - Goals is found at depth d , the shallowest goal-node depth
 - Often used in combination with DFS
- **Combines benefits of DFS and BFS**

```
function ITERATIVE_DEEPENING_SEARCH(problem)  
return a solution or failure  
  inputs: problem, a problem  
  for depth  $\leftarrow 0$  to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED_SEARCH(problem, depth)  
    if result  $\neq$  cutoff then return result
```


Iterative Deepening Search (IDS)


limit: 0



limit: 1




limit: 2



Four iterations of iterative deepening search for goal M on a binary tree, with the depth limit varying from 0 to 3.

Iterative Deepening Search (IDS)



Four iterations of iterative deepening search for goal M on a binary tree, with the depth limit varying from 0 to 3.

Iterative Deepening Search (IDS)

- **Completeness:**

- YES (no, if infinite paths)

- **Optimality:**

- YES
- Can be extended to iterative lengthening search
 - Same idea as uniform-cost search
 - Increases overhead


Iterative Deepening Search (IDS)

- **Time complexity: $O(b^d)$**

- Algorithm seems costly due to repeated generation of certain states

- Node generation:

- Level 1: d
- Level 2: $d - 1$
- ...
- Level $d-1$: 2
- Level d : 1


$$N(IDS) = db + (d - 1)b^2 + \dots + b^d$$

$$N(BFS) = b + b^2 + \dots + b^d + (b^{d+1} - b)$$

- e.g. $b=10$ and $d=5$

$$N(IDS) = 5 * 10 + 4 * 10^2 + 3 * 10^3 + 2 * 10^4 + 1 * 10^5 = 123450$$


$$N(BFS) = 10 + 10^2 + 10^3 + 10^4 + 10^5 + (10^6 - 10) = 1111100$$

Iterative Deepening Search (IDS)

- **Space complexity:** $O(bd)$
 - cf. DFS: $O(bm)$


Bidirectional Search

- **Two simultaneous searches from start and goal**
 - Motivation: $b^{d/2} + b^{d/2} < b^d$ (BFS)
 - Check whether the node belongs to the other fringe before expansion



Bidirectional Search

- **The predecessor of each node should be efficiently computable**
 - When actions are easily reversible
 - Space complexity is the most significant weakness
 - Complete and optimal if both searches are BFS



Summary of Algorithms




Criterion	Breadth-First	Uniform-cost	Depth-First	Depth-limited	Iterative deepening	Bidirectional
Complete	YES*	YES*	NO	YES, if $l \geq d$	YES	YES*
Time	b^{d+1}	$b^{C^*/e}$	b^m	b^l	b^d	$b^{d/2}$
Space	b^{d+1}	$b^{C^*/e}$	bm	bl	bd	$b^{d/2}$
Optimal	YES*	YES*	NO	NO	YES	YES

Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Repeated States

- **Failure to detect repeated states can turn a solvable problems into unsolvable ones**



Graph-Search

- **Closed list stores all expanded nodes**

- **function** GRAPH-SEARCH(*problem*, *fringe*) **return** a solution or failure
 closed ← an empty set
 fringe ← INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *fringe*)
 loop do
 if EMPTY(*fringe*) **then return** failure
 node ← REMOVE-FIRST(*fringe*)
 if GOAL-TEST[*problem*] applied to STATE[*node*] succeeds **then**
 return SOLUTION(*node*)
 if STATE[*node*] is not in *closed* **then**
 add STATE[*node*] to *closed*
 fringe ← INSERT-ALL(EXPAND(*node*, *problem*), *fringe*)

Graph-Search

- **Optimality:**

- When detecting a repeated state, i.e. matching a node on closed list, GRAPH-SEARCH discard *newly discovered* paths
 - This may result in a sub-optimal solution
 - If the new one is shorter
 - Won't happen when using uniform-cost search or BFS with constant step cost

Graph-Search

- **Time and space complexity:**
 - Proportional to the size of the state space
 - May be much smaller than $O(bd)$
 - On problem with many repeated states, GRAPH-SEARCH is much more efficient than TREE-SEARCH
 - DFS and IDS with closed list require more than linear space
 - Since all nodes are stored in closed list!

Fundamental tradeoff between **time** and **space**:
Algorithms that forget their history are doomed to repeat it!

Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Search with Partial Information

- **Previous assumption:**
 - Environment is *fully observable*
 - Environment is *deterministic*
 - Agent *knows* the effects of its actions
- **What if knowledge of states or actions is incomplete or uncertain?**

Search with Partial Information

- **Partial knowledge of states and actions**
 - Sensorless or conformant problem
 - Non-observable
 - Agent may have no idea where it is; solution (if any) is a sequence
 - Contingency problem
 - Nondeterministic and/or partially observable
 - Percepts provide new information about current state; solution is a tree or policy; often interleave search and execution
 - If uncertainty is caused by actions of **another** agent
→ adversarial problem
 - Exploration problem
 - Unknown state space
 - When states and actions of the environment are unknown

Search with Partial Information

- **Single-state problem:**


- Q: Single state, starting in #5. Solution?
- A: [Right, Suck]

- **Sensorless problem:**

- Q: Start in {1,2,3,4,5,6,7,8}
e.g. Right goes to {2,4,6,8}.
Solution?

- **Contingency:**

- Q: Start in {1,3}.
Assume Murphy's law, Suck
can dirty a clean carpet and
local sensing: [*location, dirt*]
only. Solution?



Sensorless Problems


- **e.g. Vacuum world**

- Q: Start in {1,2,3,4,5,6,7,8}
e.g. Right goes to {2,4,6,8}.
Solution?

- A: [Right, Suck, Left, Suck]

- **When the world is not fully observable:**


- reason about a set of states that **might be reached**
→ belief state
 - If fully observable, each belief state contains only one physical state




Sensorless Problems

- **Search space of belief states**

- Search the space of belief states rather than physical states
- Solution = a path to a belief state, *all* of whose members are goal states



Belief-state of Vacuum World




Contingency Problems

- **e.g. Vacuum world**

- Start in {1,3},
Murphy's law: Suck can dirty a clean carpet

Local sensing: dirt, location

- Percept = [L,Dirty] = {1,3}
- [Suck] = {5,7}
- [Right] = {6,8}
- [Suck] in {6}={8} (Success)
- BUT [Suck] in {8} may return {6} (Murphy's law)
- Solution??
 - Belief-state: no fixed action sequence guarantees solution



Contingency Problems

- **Relax requirement**

- If don't insist on fixed action sequence, then a solution:

[Suck, Right, if [R,dirty] then Suck]

- Select actions based on contingencies arising during execution
 - Most real-world problems are contingency problems, because exact prediction is impossible

Outline

- **Problem-solving agents**
- **Problem formulation**
- **Search for solutions**
- **Uninformed search strategies**
- **Avoiding repeated states**
- **Search with partial information**
- **Problem difficulty**

Reasons of Difficult Problems

- **Huge search space**
 - Exhaustive search for the best answer is forbidden
- **Complicated**
 - Result from the simplified models is useless
- **Dynamic**
 - The evaluation function is noisy or varies with time
 - Thereby needs not just a single solution but a series of ones
- **Constrained**
 - Even on a feasible answer is difficult, let alone the optimum
- **Multiple objectives**
 - Conflict between objectives

Size of Search Space

- **Satisfiability problem (SAT)**

- Given a logic formula in CNF, the problem is to determine if the formula is satisfiable

- E.g. 100-variable SAT:

$$f(\vec{x}) = (x_1 \vee x_5 \vee \neg x_3) \wedge (x_6) \wedge (\neg x_4 \vee x_2)$$

- For a n-variable SAT, the size of search space is 2^n

- E.g. $|S| = 2^{100} \approx 10^{30}$

For a computer that can test 1000 strings/sec, from Big Bang to now (15 billion years), we'd have examined $< 1\%$

- **Evaluation function** is another issue

- Evaluation function guides us the quality of solution
 - What if SAT?

Size of Search Space

- **Traveling salesman problem (TSP)**
 - Find the shortest path to visit each city exactly
 - Symmetric and asymmetric
 - The size of search space is $n!/2n = (n-1)!/2$
 - E.g. 50-city TSP: $|S| = 49!/2 \approx 10^{62}$
 - 100-city TSP: $|S| \approx 4.67 \times 10^{155}$
 - 500-city TSP: $|S| \approx 1.22 \times 10^{1131}$
 - 3000-city TSP: $|S| \approx 6.91 \times 10^{9126}$



TSP instance with 24,978 cities

Size of Search Space

- **Nonlinear problem (NLP)**

- Find the global maximum of function G2

$$G_2(\vec{x}) = \left| \frac{\sum_{i=1}^n \cos^4(x_i) - 2 \prod_{i=1}^n \cos^2(x_i)}{\sqrt{\sum_{i=1}^n i x_i^2}} \right|$$

s.t. $\prod_{i=1}^n x_i \geq 0.75, \sum_{i=1}^n x_i \leq 7.5n$

- Nonlinear
- Constraints: 1 non-linear, 1 linear
- Search space
 - Infinitely large due to real
 - Depends on resolution
 - E.g. precision guaranteed 6 digits
each dimension has 10^7 values
search space = 10^{7n}

Modeling the Problem

- **Solving a problem means we are in reality only find the solution to a **model** of the problem**
- **Process**
 - Create a model of the problem
 - Use the model to generate a solution



Modeling the Problem

- **E.g. SAT, TSP, NLP are canonical models**
 - SAT: model checking
 - TSP: die bonding, factory scheduling
 - NLP: function optimization
- **However, there is always a compromise on precision between model and solution**

Modeling the Problem

- Example

- Suppose a company has n warehouses that store paper supplies in reams
- These supplies are to be delivered to k distribution centers
- The transportation cost between warehouse i and distribution center j is f_{ij} :

$$f_{ij}(x) = \begin{cases} 0 & x = 0 \\ 4 + 3.33x & 0 < x \leq 3 \\ 19.5 & 3 < x \leq 6 \\ 0.5 + 10\sqrt{x} & 6 < x \end{cases}$$

- The problem is to minimize the transportation cost of all:

$$\min \sum_{i=1}^n \sum_{j=1}^k f_{ij}(x_{ij}) \quad \text{s. t.} \quad \sum_{j=1}^k x_{ij} \leq \text{warehouse}(i), \sum_{i=1}^n x_{ij} \geq \text{distribution}(j), x_{ij} \geq 0$$

Modeling the Problem

- **Example** (cont'd)

- The transportation cost serves as an evaluation function
 - Discontinuous \rightarrow cannot use gradient-based methods
- Two options
 - Simplify the model so that traditional optimizers might return better answers:
Problem \rightarrow Model_a \rightarrow Solution_p(Model_a)
 - Keep the model but use non-traditional approach to find a near-optimum solution:
Problem \rightarrow Model_p \rightarrow Solution_a(Model_p)
- Difficult to obtain precise solution since we have to approximate either the *model* or the *solution*

Change over Time

- **Problem (environment) changes**
 - Before you model it
 - While you are deriving a solution
 - After you execute your solution
- **E.g.**
 - TSP: the cost changes over time
 - Game: your competitors are trying to defeat your solution

Constraints

- During search, we need to move from one solution to the next:
 - Feasible?
 - Improved?
- **Hard** constraints
 - Must be absolutely satisfied in order to have a feasible solution
- **Soft** constraints
 - We hope to accomplish but are NOT mandatory

Any solution that meets the hard constraints is **feasible**, but not necessary **optimal** in light of soft constraints.

Constraints

- **Deal with infeasible solutions**
 - Discard
 - Diversity
 - Repair
 - Diversity
 - Computational cost
 - Prevent
 - Devise variation operators that never corrupt a feasible solution into an infeasible
 - Not always achievable

Multiple Objectives

- **Presence of multiple conflicting objectives**
 - Buying a car: speed vs. price vs. reliability
 - Engineering design: lightness vs. strength
 - Factors are conflicting and the final decision bears a compromise, a trade-off
- **Two part problem**
 - Finding set of good solutions
 - Choice of the best for particular application

Multiple Objectives

- **Methods**

- Aggregation (scalarization)
 - Use weights to integrate into a single objective
 - Disadvantages?
- Evolutionary algorithms
 - Good at solving MOP
 - Returns a set of solutions
 - NSGA2, PEAS, MOEA, etc.

